

Large-Scale Automatic K-Means Clustering for Heterogeneous Many-Core Supercomputer

Teng Yu, Wenlai Zhao, Pan Liu, Vladimir Janjic, Xiaohan Yan, Shicai Wang Haohuan Fu, Guangwen Yang, John Thomson

汇报人：胡畅

College of Computer Science and Electronic Engineering

October 5, 2022

1 Introduction

2 Methods

- DataFlow Partition
- DataFlow and Centroids Partition
- DataFlow and Centroids and Dimensions Partition
- Determining the optimal k

3 Experiment Results

4 Summary

Outline

- 1 Introduction
- 2 Methods
- 3 Experiment Results
- 4 Summary

K-Means Definition

Formalized, given n samples, $\mathcal{X}^d = \{x_i^d | x_i^d \in \mathcal{R}^d, i \in \{1, \dots, n\}\}$, where $x_i^d = (x_{i1}, x_{i2}, \dots, x_{id})$, We aim to find k d-dimensional centroids $\mathcal{C}^d = \{c_j^d | c_j^d \in \mathcal{R}^d\}, j \in \{1, \dots, k\}$ to minimize the object $\mathcal{O}(\mathcal{C})$:

$$\mathcal{O}(\mathcal{C}) = \frac{1}{n} \sum_{i=1}^n \text{dis}(x_i^d, c_{a_i}^d)$$

$$a_i = \operatorname{argmin}_{j \in \{1, \dots, k\}} \text{dis}(x_i^d, c_j^d)$$

$$1. : a(i) = \operatorname{arg min}_{j \in \{1 \dots k\}} \text{dis}(x_i^d, c_j^d) \text{ (Assign)}$$

$$2. : c_j^d = \frac{\sum_{\operatorname{arg} a(i)=j} x_i^d}{|\operatorname{arg} a(i) = j|} \text{ (Update)}$$

Figure: KMeans algorithm steps

Example of K-means




Figure: KMeans example

Why parallel K-Means?

Finding the optimal solution for a general k-means problem is known to be NP-hard:

- 1 number of centroids (k)
- 2 number of dimensions (d)
- 3 proper hyper-parameters, such as the targeted number of centroids (k)

- large-scale clustering problems with up to 196,608 dimensions and over 160,000 targeting centroids

Introduction

- large-scale clustering problems with up to 196,608 dimensions and over 160,000 targeting centroids
- automatic hyper-parameter determination

- large-scale clustering problems with up to 196,608 dimensions and over 160,000 targeting centroids
- automatic hyper-parameter determination
- achieve high performance and scalability for problems

- large-scale clustering problems with up to 196,608 dimensions and over 160,000 targeting centroids
- automatic hyper-parameter determination
- achieve high performance and scalability for problems
- support clustering without sufficient prior knowledge

1 Introduction

2 Methods

- DataFlow Partition
- DataFlow and Centroids Partition
- DataFlow and Centroids and Dimensions Partition
- Determining the optimal k

3 Experiment Results

4 Summary

Sunway TaihuLight

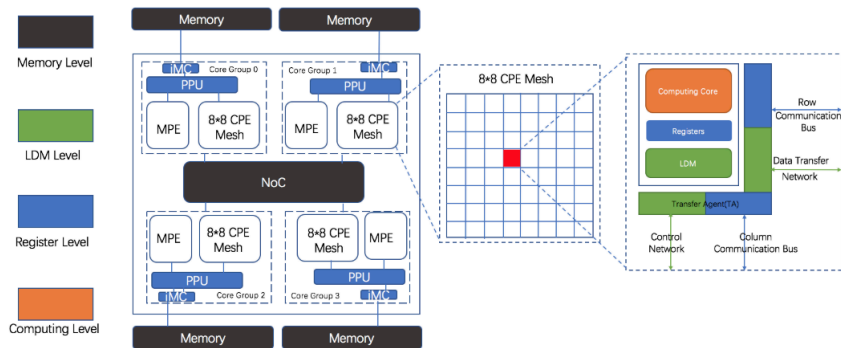


Figure: The general architecture of the SW26010 many-core processor

Three level parallelism

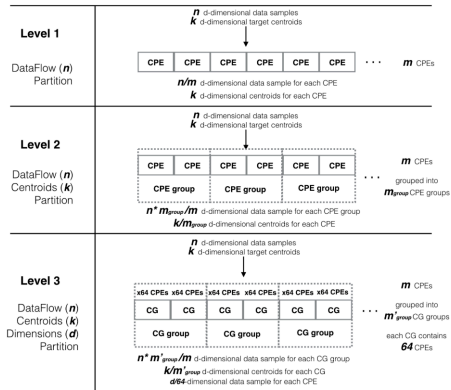


Fig. 2. Three-level k -means design for data partition and parallelism on Sunway architecture

Figure: Three-level k-means design for data partition and parallelism on Sunway architecture

3.1 Level 1 - DataFlow Partition

Algorithm 1 Basic Parallel k -means

```
1: INPUT: Input dataset  $\mathcal{X} = \{x_i | x_i \in R^d, i \in [1, n]\}$ , and  
   initial centroid set  $\mathcal{C} = \{c_j | c_j \in R^d, j \in [1, k]\}$   
2:  $P_l \xleftarrow{\text{load}} \mathcal{C}, l \in \{1 \dots m\}$   
3: repeat  
4:   // Parallel execution on all CPEs:  
5:   for  $l = 1$  to  $m$  do  
6:     Init a local centroids set  $\mathcal{C}^l = \{c_j^l | c_j^l = \mathbf{0}, j \in [1, k]\}$   
7:     Init a local counter  $count^l = \{count_j^l | count_j^l = 0, j \in [1, k]\}$   
8:     for  $i = (1 + (l - 1) * \frac{n}{m})$  to  $(l * \frac{n}{m})$  do  
9:        $P_l \xleftarrow{\text{load}} x_i$   
10:       $a(i) = \arg \min_{j \in \{1 \dots k\}} dis(x_i, c_j)$   
11:       $c_{a(i)}^l = c_{a(i)}^l + x_i$   
12:       $count_{a(i)}^l = count_{a(i)}^l + 1$   
13:    end for  
14:    for  $j = 1$  to  $k$  do  
15:      AllReduce  $c_j^l$  and  $count_j^l$   
16:       $c_j^l = \frac{c_j^l}{count_j^l}$   
17:    end for  
18:  end for  
19: until  $\mathcal{C}^l == \mathcal{C}$   
20: OUTPUT:  $\mathcal{C}$ 
```

Figure: Level 1 - DataFlow Partition

DataFlow Partition

```
1
2
3 #include <omp.h>
4 #include <stdio.h>
5
6 #define len 10000
7 int data[len];
8 int threads[10];
9
10 int main() {
11
12     for(int i = 0; i < len; i++) {
13         data[i] = 1;
14     }
15
16     #pragma omp parallel for num_threads(10)
17     for(int i = 0; i < len; i++) {
18         printf("%d i = %d\n", (int) omp_get_thread_num(), i);
19         threads[(int) omp_get_thread_num()] += data[i];
20     }
21     int sum = 0;
22     for(int i = 0; i < 10; i++) {
23         printf("threads[%d] = %d\n", i, threads[i]);
24         sum += threads[i];
25     }
26     printf("sum = %d\n", sum);
27     return 0;
28 }
```

Figure: Openmp code for parallel

DataFlow and Centroids Partition

Algorithm 2 Parallel *k-means* for *k*-scale

```
1: INPUT: Input dataset  $\mathcal{X} = \{x_i | x_i \in R^d, i \in [1, n]\}$ , and  
   initial centroid set  $\mathcal{C} = \{c_j | c_j \in R^d, j \in [1, k]\}$   
2:  $\underline{P}_j \xleftarrow{\text{load}} c_j$   $j \in (1 + \text{mod}(\frac{l-1}{m_{\text{group}}}) * \frac{k}{m_{\text{group}}},$   
    $\text{mod}(\frac{l-1}{m_{\text{group}}}) + 1) * \frac{k}{m_{\text{group}}}$   
3: repeat  
4:   // Parallel execution on each CPE group  $\{P\}_{l'}$ :  
5:   for  $l' = 1$  to  $\frac{m}{m_{\text{group}}}$  do  
6:     Init a local centroids set  $\mathcal{C}^{l'}$  and counter  $\text{count}^{l'}$   
7:     for  $i = (1 + (l' - 1) \frac{n * m_{\text{group}}}{m})$  to  $(l' \frac{n * m_{\text{group}}}{m})$  do  
8:        $\{P\}_{l'} \xleftarrow{\text{load}} x_i$   
9:        $a(i)^{l'} = \arg \min_j \text{dis}(x_i, c_j)$   
10:       $a(i) = \min. a(i)^{l'}$   
11:       $c'_{a(i)} = c^{l'}_{a(i)} + x_i$   
12:       $\text{count}^{l'}_{a(i)} = \text{count}^{l'}_{a(i)} + 1$   
13:    end for  
14:    for  $j = (1 + \text{mod}(\frac{l-1}{m_{\text{group}}}) * \frac{k}{m_{\text{group}}})$  to  
       $((\text{mod}(\frac{l-1}{m_{\text{group}}}) + 1) * \frac{k}{m_{\text{group}}})$  do  
15:      AllReduce  $c'^{l'}_j$  and  $\text{count}^{l'}_j$   
16:       $c'^{l'}_j = \frac{c'^{l'}_j}{\text{count}^{l'}_j}$   
17:    end for  
18:  end for  
19: until  $\cup \mathcal{C}^{l'} == \mathcal{C}$   
20: OUTPUT:  $\mathcal{C}$ 
```

Figure: Level 2 - DataFlow and Centroids Partition

DataFlow and Centroids and Dimensions Partition

Algorithm 3 Parallel k -means for k -scale and d -scale

```
1: INPUT: Input dataset  $\mathcal{X} = \{x_i | x_i \in R^d, i \in [1, n]\}$ , and  
   initial centroid set  $\mathcal{C} = \{c_j | c_j \in R^d, j \in [1, k]\}$   
2:  $CG_{l''} \xleftarrow{\text{load}} c_j^d, l'' \in \{1 \dots \frac{m}{64}\}, j \in (1 + \text{mod}(\frac{l''-1}{m'_{group}}) * \frac{k}{m'_{group}}, (\text{mod}(\frac{l''-1}{m'_{group}}) + 1) * \frac{k}{m'_{group}})$   
3: repeat  
4:   // Parallel execution on each CG group  $\{CG\}_{l''}$ :  
5:   for  $l'' = 1$  to  $\frac{m}{64}$  do  
6:     Init a local centroids set  $\mathcal{C}^{l''}$  and counter  $count^{l''}$   
7:     for  $i = (1 + (l'' - 1) * \frac{n * m'_{group}}{m})$  to  $(l'' * \frac{n * m'_{group}}{m})$  do  
8:       for  $u = (1 + \text{mod}(\frac{l-1}{64}) * \frac{d}{64})$  to  $(\text{mod}(\frac{l-1}{64}) + 1) * \frac{d}{64}$  do  
9:          $CG_{l''} \leftarrow x_i (P_i \leftarrow x_i^u)$   
10:      end for  
11:       $a(i)' = \arg \min_j \text{dis}(x_i, c_j)$   
12:       $a(i) = \min. a(i)'$   
13:       $c_{a(i)}^{l''} = c_{a(i)}^{l''} + x_i$   
14:       $count_{a(i)}^{l''} = count_{a(i)}^{l''} + 1$   
15:    end for  
16:    for  $j = (1 + \text{mod}(\frac{l''-1}{m'_{group}}) * \frac{k}{m'_{group}})$  to  
       $((\text{mod}(\frac{l''-1}{m'_{group}}) + 1) * \frac{k}{m'_{group}})$  do  
17:      AllReduce  $c_j^{l''}$  and  $count_j^{l''}$   
18:       $c_j^{l''} = \frac{c_j^{l''}}{count_j^{l''}}$   
19:    end for  
20:  end for  
21: until  $\cup \mathcal{C}^{l''} == \mathcal{C}$   
22: OUTPUT:  $\mathcal{C}$ 
```

Figure: Level 3 - DataFlow and Centroids and Dimensions Partition

Determining the optimal

$$r(k) = \inf\{t : \exists y_1, \dots, y_k \text{ in } R^d, \mathcal{X}^d \subseteq \bigcup_{1 \leq s \leq k} B(y_s, t)\},$$

$$\Delta r'(k) = r'(k) - r'(k+1),$$

Figure: Formula of determining the optimal k




Figure: The evaluation function $r'(k)$ to determine the optimal k value.

Outline

- 1 Introduction
- 2 Methods
- 3 Experiment Results**
- 4 Summary

Experiment Results

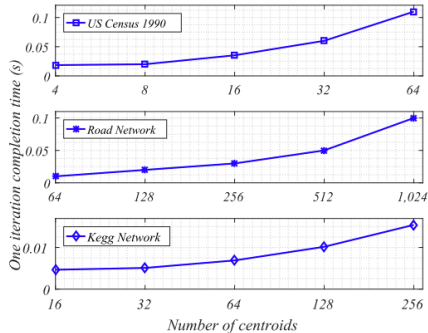


Figure: Level 1 - dataflow partition




Figure: Level 2 - dataflow and centroids partition

Experiment Results




Figure: Level 3 - dataflow, centroids and data-sample partition

Experiment Results




Figure: varying d with 2,000 centroids and 1,265,723 data samples tested on 128 nodes

Experiment Results




Figure: varying number of nodes used with a xed 4,096 dimension, 2,000 centroids and 1,265,723 data samples

TABLE 3
Execution time comparison with other architectures

Approaches	Hardware Resources	n	k	d	Execution time per iteration (sec.)	Execution time per iteration by Sunway Taihu-Light (sec.)	Max. Speedup
Rossbach, et al [35]	10x NVIDIA Tesla K20M + 20x Intel Xeon E5-2620	1.0E9	120	40	49.4	0.468635 (128 nodes)	105x
Bhimani, et al [4]	NVIDIA Tesla K20M	1.4E6	240	5	1.77	0.025336 (4 nodes)	70x
Jin, et al [26]	NVIDIA Tesla K20c	1.4E5	500	90	5.407	0.110191 (1 node)	49x
Li, et al [29]	Xilinx ZC706	2.1E6	4	4	0.0085	0.002839 (1 node)	3x
Ding, et al [15]	Intel i7-3770K	2.5E6	10,000	68	75.976	2.424517 (16 nodes)	31x

Figure: Execution time comparison with other architectures

Outline

- 1 Introduction
- 2 Methods
- 3 Experiment Results
- 4 Summary**

Contributions of this paper:

- ① Level 1 - DataFlow Partition: Store a whole sample and k centroids on single-CPE
- ② Level 2 - DataFlow and Centroids Partition: Store a whole sample on single-CPE whilst k centroids on multi-CPE
- ③ Level 3 - DataFlow, Centroids and Dimensions Partition: Store a whole sample on multi-CPE whilst k centroids on Multi-CG and d dimensions on Multi-CPE
- ④ The proposed auto-clustering solution is a significant attempt to support AutoML on a supercomputer system

Thanks for your attention!